layers 1–3: harness, CLAUDE.md, AGENTS.md/SOUL.md — identical to every other skill, see diagram above

4 · SKILL BODY

.claude/skills/deep-research/SKILL.md
the 7 phases
.claude/skills/deep-research/references/report.md
exact protocol, depth-preset defaults,
the report template, the peer-review reconciliation rule — pulled in because
SKILL.md cites it by name rather than
repeating it

5 · CROSS-GENRE STANDARDS
same as every other skill

docs/WRITING-STANDARDS.md
assets/style/acronyms.toml

6 · RUN STATE
grows phase by phase — most of what makes a deep-research prompt large by the time Phase 7 runs

content/dossiers/deep-research-<slug>/scope.md
content/dossiers/deep-research-<slug>/evidence.md
every interviewer's kept claims (Phase 2) and
every writer's added sources (Phase 5),
transcribed by the orchestrator itself —
a subagent never writes here
content/dossiers/deep-research-<slug>/rejected.md
every discarded citekey, with the query and why
content/dossiers/deep-research-<slug>/sections.md
the outline → citekey plan (Phase 4), reconciled
to what the report actually cites (Phase 7e)
subagent packets returned from Phases 2, 5 and 7
read once, transcribed above, then gone

7 · THIS TURN

topic · reader · scope · depth preset
(quick / standard / deep) · any steering